

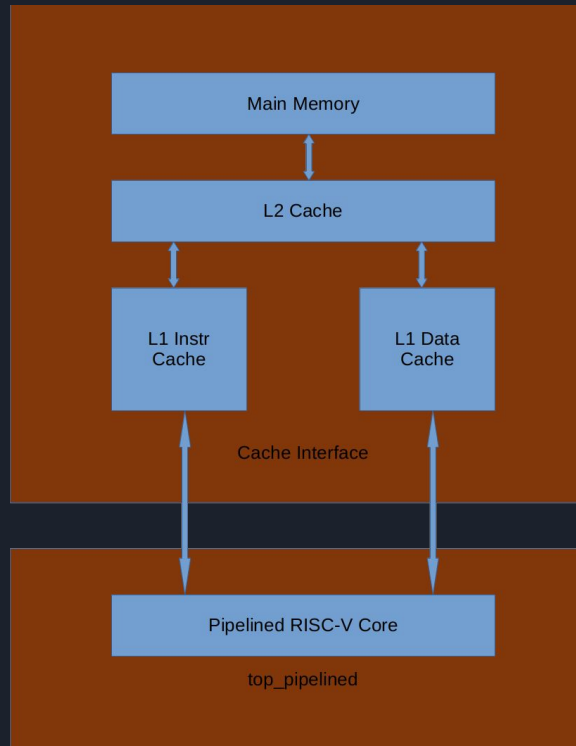


Multi-core Processor on FPGA

Lasya Balachandran and Sanjay Seshan
6.1920[6.175] Final Project

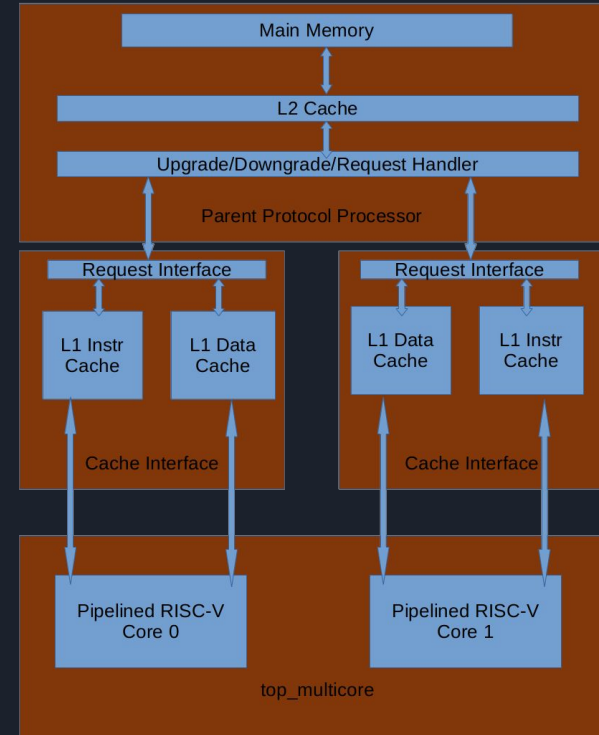
Connecting Cache to the Processor

- Converted cache to store words
 - Used 4 offset bits to correspond to 16 words per line
- Used two BRAMs for the cache
 - One with tag and valid
 - One byte enabled for data
- Created two caches (data and instruction)
- Requests for instructions and data are tagged in a queue when sending to L2 to return to the cache in FIFO order
- 2 level cache hierarchy connects to main memory
 - Built an interface between processor and cache
- Main memory and L2 cache store lines
- Evaluated using modified Beveren tests



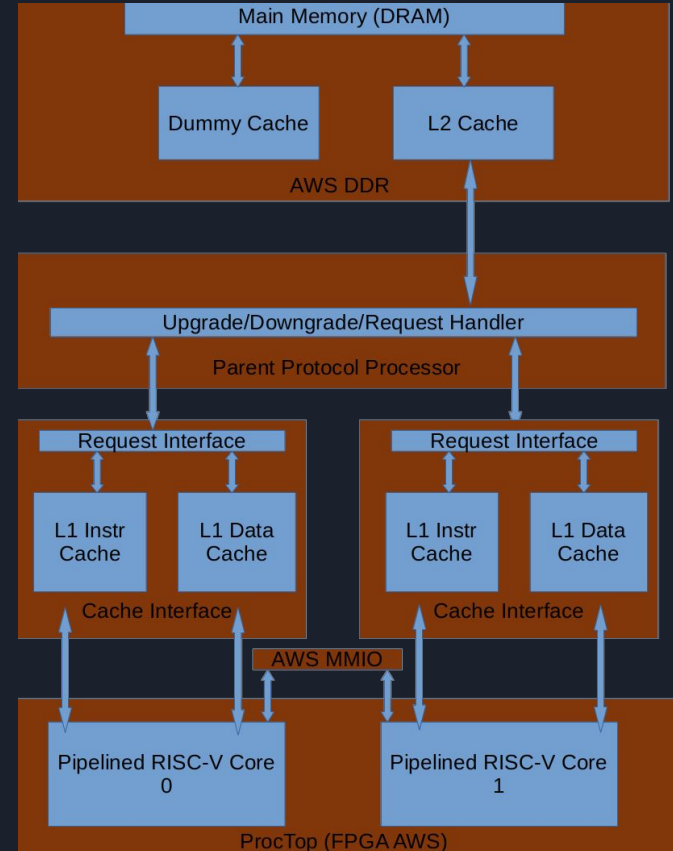
Implementing Multi-Core Processor

- Created a Parent Protocol Processor
 - Routes upgrades and downgrades between cores
 - Abstracted the L2 cache and main memory into the PPP
 - PPP handles requests for evictions and loads into higher cache levels
 - Shared L2 cache stores writes for accessing by other cores
 - Tag core requests to memory to return FIFO order
- System is modular
- Two cores are instantiated
 - First starts at pc=0
 - Second starts at pc=4096
- Each core gets its own set of data and instruction caches



Running on FPGA

- We modified our code to run within the AWS DDR system
- Restructured the cache hierarchy a bit to match
- Had to use a dummy cache to use a single L2 cache with no separate instruction L2 cache
- MMIO was also abstracted into the AWS module and was shared by the cores by using a FIFO tag on requests
- Only managed to get a few tests working in time but results are promising



Results – Multicore

- Standard tests run same code on both cores which all work
- New tests ensures coherence between core's caches
- Created a distributed matrix multiply to display working shared data and speed up from two threads
- Each core produced half of the rows
- Thread1 notifies Thread0 when it finishes
- **Single core ran with 5.47 million cycles** (simulator 2m 43s)
- **Dual core ran with 2.67 million cycles** (simulator 2m 22s)
 - Roughly 2x cycle speed up
 - Simulator has overhead and does not really utilize multi core processing on the host

```
1 void program_thread0(){
2     int sum;
3     for (int i = 0; i < 8; i++)
4     {
5         for (int j = 0; j < 16; j++)
6         {
7             sum = 0;
8             for (int k = 0; k < 16; k++)
9             {
10                 sum += multiply(a[i][k], b[k][j]);
11             }
12             c[i][j] = sum;
13             putchar(sum);
14         }
15     }
16
17     while (flag == 0); // Wait until thread1 produced the value
18     if (arrEquals(expected, c))
19     {
20         exit(0);
21     }
22     else
23     {
24         exit(1);
25     }
26 }
27 }
```

```
30 void program_thread1(){
31     int sum;
32     for (int i = 8; i < 16; i++)
33     {
34         for (int j = 0; j < 16; j++)
35         {
36             sum = 0;
37             for (int k = 0; k < 16; k++)
38             {
39                 sum += multiply(a[i][k], b[k][j]);
40             }
41             c[i][j] = sum;
42             putchar(sum);
43         }
44     }
45     // mostly since the C compiler is too smart to do flag = 1 or something
46     int l = 0;
47     for (int i = 0; i < 4; i++){
48         l += i;
49     }
50     flag = 1;
51     while(1);
52 }
```

Results – FPGA

- Many of the single core tests worked on the FPGA
- Remaining tests are a work in progress
- Clock speeds are shown at the right and are promising

Clock Summary			
Clock	Waveform(ns)	Period(ns)	Frequency(MHz)
CLK_300M_DIMM0_DP	{0.000 1.666}	3.332	300.120
mmcm_clkout0	{0.000 1.874}	3.749	266.773
pll_clk[0]	{0.000 0.234}	0.469	2134.187
pll_clk[0]_DIV	{0.000 1.874}	3.749	266.773
pll_clk[1]	{0.000 0.234}	0.469	2134.187
pll_clk[1]_DIV	{0.000 1.874}	3.749	266.773
pll_clk[2]	{0.000 0.234}	0.469	2134.187
pll_clk[2]_DIV	{0.000 1.874}	3.749	266.773
mmcm_clkout6	{0.000 3.749}	7.497	133.387
refclk_100	{0.000 5.000}	10.000	100.000
qplloutclk_out[3]	{0.000 0.100}	0.200	5000.001
txoutclk_out[15]	{0.000 1.000}	2.000	500.000
clk_core	{0.000 2.000}	4.000	250.000
clk_main_a0	{0.000 4.000}	8.000	125.000
tck	{0.000 16.000}	32.000	31.250



Challenges

- Handling half words and bytes took some restructuring
- Many revisions and debugging to ensure PPP handles correct cache coherency
- Some issues turning words to lines in mem.vmh – needed padding of zeroes
- Unexpected loading values in the cache – had to zero out the valid/dirty bits

Thanks
for
listening!